

# STAT 9100 HW1

Mengyan Jing

2026-02-09

## Instructions

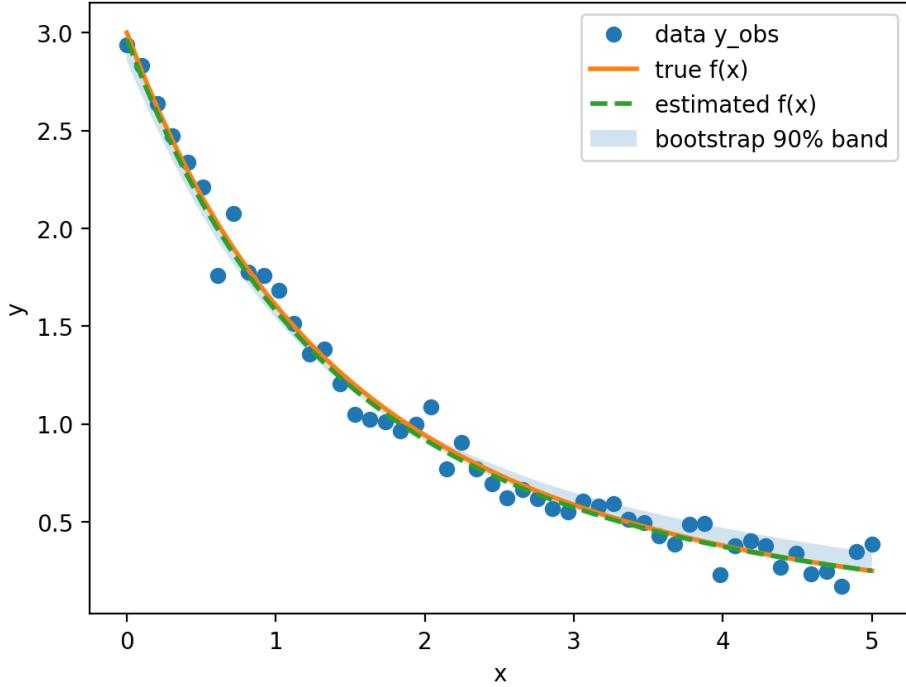
This homework explores neural estimation and conformal inference for a nonlinear regression model.

### Problem 1

We trained a neural Bayes estimator to map the standardized observation vector  $y \in \mathbb{R}^{50}$  to the parameter vector  $\theta = (\beta_1, \gamma_1, \beta_2, \gamma_2)$  using simulated training pairs. To quantify uncertainty for the single observed curve, we used a residual bootstrap with 100 resamples.

Table 1: Neural Bayes estimate and bootstrap uncertainty (B=100).

Parameter	Estimate	BootMean	BootSD	CI2.5	CI97.5
beta1	1.522	1.530	0.076	1.375	1.669
gamma1	0.369	0.347	0.032	0.279	0.395
beta2	1.443	1.408	0.090	1.248	1.578
gamma2	1.010	1.058	0.089	0.908	1.244



## Problem 2

### Settings (Neural Likelihood)

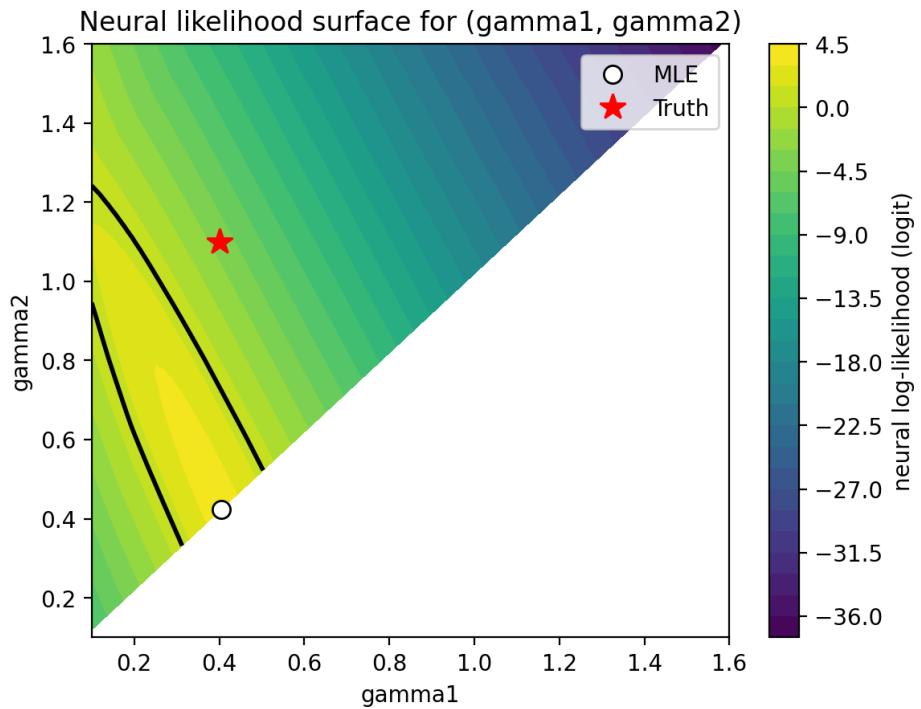
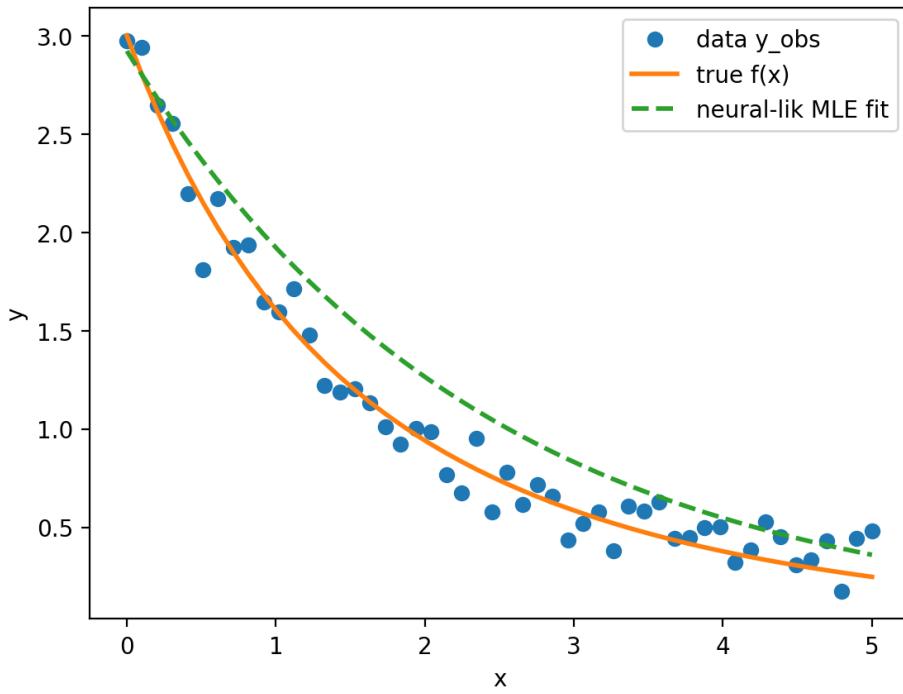
- Model:  $y(x) = \beta_1 e^{-\gamma_1 x} + \beta_2 e^{-\gamma_2 x} + \epsilon$ ,  $\epsilon \sim N(0, \sigma^2)$ ,  $\sigma = 0.1$ .
- Design points:  $x$  is a fixed grid of 50 points on  $[0, 5]$ .
- Prior (training simulations):  $\beta_1, \beta_2 \sim \text{Unif}(0.5, 2.5)$ ,  $\gamma_1, \gamma_2 \sim \text{Unif}(0.1, 1.6)$  with identifiability constraint  $\gamma_1 < \gamma_2$ .
- Training data:  $N = 60,000$  simulated pairs.
- Classifier architecture: fully-connected network  $(54 \rightarrow 128 \rightarrow 128 \rightarrow 1)$  with ReLU.
- Training: Adam (lr=1e-3), batch size 512, epochs 8; response curves standardized using global mean/std from the simulated training set.
- MLE search: multi-start gradient ascent over a constrained parameterization (via sigmoid map to bounds and sorting to enforce  $\gamma_1 < \gamma_2$ ); warm-started using Problem 1 estimate.

Table 2: Problem 2: Neural-likelihood approximate MLE and truth.

Parameter	Estimate	Truth
beta1	0.678	1.8
gamma1	0.405	0.4
beta2	2.245	1.2
gamma2	0.423	1.1

Table 3: Problem 2: Likelihood-ratio diagnostic at the truth.

best_logit	LR_at_truth	LR95_thresh_df2
3.854	3.263	5.99



## Conclusion

We trained a neural likelihood classifier to distinguish *paired* samples  $(y, \theta) \sim p(y | \theta)p(\theta)$  from *mismatched* samples obtained by permuting  $\theta$ . The classifier logit evaluated at the observed curve  $y_{\text{obs}}$  provides an approximation to the log-likelihood surface  $\ell(\theta; y_{\text{obs}})$  up to an additive constant. Maximizing this neural log-likelihood over a constrained parameterization (positivity and the identifiability constraint  $\gamma_1 < \gamma_2$ ) yielded an approximate MLE

$$\hat{\theta}_{\text{MLE}} = (\hat{\beta}_1, \hat{\gamma}_1, \hat{\beta}_2, \hat{\gamma}_2) = (0.678, 0.405, 2.245, 0.423),$$

with best logit value 3.854. The fitted response captures the overall exponential decay pattern in the data, but deviates from the true response for larger  $x$ , suggesting that multiple parameter combinations can produce similar fitted curves in this two-exponential mixture model.

The two-dimensional neural likelihood surface over  $(\gamma_1, \gamma_2)$  exhibits an elongated ridge concentrated near  $\gamma_1 \approx \gamma_2$ , indicating weak identifiability of the two-exponential decomposition. This likelihood geometry explains why optimization can favor solutions with nearly equal decay rates and redistributed weights, even when the true decay rates are more separated. Uncertainty for  $(\gamma_1, \gamma_2)$  was quantified using a likelihood-ratio (LR) region based on the grid-evaluated surface. The LR statistic at the true parameters is

$$\text{LR}(\theta_{\text{true}}) = 3.263 < \chi^2_2(0.95) = 5.99,$$

so the true  $(\gamma_1, \gamma_2)$  lies inside the 95% LR region, consistent with the contour overlay in the likelihood surface plot.

## Problem 3

## Problem 4